Summary – Learning Through Narratives: Sample-Efficient Intelligence via Episodic Compression to Latent Variables

Executive Summary

This work proposes an alternative to embodied learning for achieving sample-efficient artificial intelligence. While LeCun (2022) argues that physical interaction is necessary for rapid learning, we demonstrate that narrative-based learning – how humans seem to acquire substantial knowledge – provides a practical path forward using existing language models. Through episodic compression, experiences compress into salient episodes that generate generalizations via pattern detection, enabling prediction from minimal examples without requiring neural network training.

The Sample Efficiency Challenge

Modern machine learning faces a fundamental problem: systems require millions of training examples to achieve competence that humans develop from handfuls of experiences. A person learns to drive in approximately 20 hours; children grasp physics and social rules from limited exposure. Meanwhile, state-of-the-art ML systems demand vast datasets.

LeCun's Joint Embedding Predictive Architecture (JEPA) addresses this through embodied physical interaction, learning representations via gradient descent on sensorimotor data. This requires training hierarchical networks from scratch over weeks on specialized hardware.

The Narrative Learning Alternative

Consider how humans acquire knowledge. Therapists understand psychological conditions through case studies and literature, not personal experience of every disorder. Children learn dangers primarily through warnings and stories. Scientists build on accumulated literature rather than personally validating every finding. This suggests narrative learning may suffice for many domains.

LeCun’s work acknowledges humans learn extensively through observation – which differs from narratives only by requiring temporal and spatial proximity. Narratives enable learning from events across time and space without presence, potentially offering broader accessibility than embodied interaction.

Large Language Models as Cognitive Substrate

LLMs provide suitable infrastructure for narrative-based learning through three key properties:

Natural language tasking: Rather than implementing pattern-matching algorithms from scratch, you can instruct the system to "find patterns."

Attention mechanisms: Transformer architecture performs parallel pattern-matching across contexts.

Baseline reasoning: Pre-training on massive text corpora provides understanding of natural language, baseline physics and causation knowledge, and ability to follow instructions and perform consistency checking.

Critically, we don't ask the LLM to "learn from narratives" in the sense of having seen stories during training. The LLM performs computational operations – namely attention and reasoning – that our architecture orchestrates into a learning system.

The Episodic Compression Mechanism

The architecture implements progressive compression through three stages:

Experiences are continuous streams of interaction from any source: direct sensory input, conversational exchanges, narrative exposure, or simulated environments. Most experiences fade without retention.

Episodes are compressed summaries of experiences marked as salient. Salience derives from novelty – surprising events or prediction failures. Episodes preserve causal structure while discarding irrelevant details, naturally taking narrative format: temporal sequences with explicit causal relationships and salience markers.

Generalizations are patterns extracted when multiple episodes share structure. Pattern detection identifies commonalities across episodes, forming stable predictions. We propose a threshold mechanism: approximately 3-5 related episodes with sufficient salience generate a stable generalization, explaining both rapid learning from few examples and resistance to spurious patterns from single unusual events.

Learning happens through salience detection marking experiences for retention, episode formation compressing salient experiences, pattern extraction during consolidation, generalization formation when patterns cross threshold, and confidence tracking that strengthens or weakens with evidence.

This differs fundamentally from neural network training. Episodes and generalizations serve as learned latent variables created through salience-driven compression rather than gradient descent, enabling sample-efficient learning without parameter updates.

Experimental Validation

Testing episodic compression faces a challenge: LLMs have extensive pre-training on physics, causation, and social dynamics. We require domains where the model has no prior exposure.

We designed the Shimmer Valleys, an artificial environment with internally consistent but unprecedented causal rules. Entities include Globs (rolling, color-changing), Whisps (smoke-like, sound-responsive), Resonators (crystalline, emit tones), Flutter seeds (paper-like, fold/unfold), Shade pools (dark liquid with memory properties), and Living glass (color-changing ground surface).

Sample rules: contact transfers color through chains, sound frequencies trigger entity responses, celestial positions trigger world changes, whisp choruses transform surfaces, some transformations are irreversible, physical patterns transfer through shade pool contact. These rules maintain internal consistency but correspond to no physical laws or pre-training knowledge.

Protocol

The raw Shimmer Valleys experience compressed to a World Model delivered to the learning system:

- 9 generalizations about Shimmer Valleys causation
- 4 retained episode memories of specific events
- No access to raw experiences or full event histories
- Total state: approximately 500 words

Five novel test scenarios required combining multiple generalizations:

T1: Chrome flutter seed (from mirror surface) touches resonator. The system predicted the resonator would adopt reflective quality, creating a "mirror resonator" with visible sound patterns, combining contact transfer with information persistence. Confidence: Medium.

T2: Glowing glob (from moon triangle + resonator) enters shade pool. The system predicted the pool would absorb the glow, become luminescent, and transfer temporary bioluminescence back to the glob, applying shade pool memory/transfer properties to a novel property type. Confidence: Medium-High.

T3: Multiple whisps form chorus near silence-sphere resonator. The system predicted whisps would form an "anti-chorus" creating zones of stillness, recognizing this as inversion of known whisp chorus effects through sophisticated analogical reasoning. Confidence: Low-Medium.

T4: Phase glob encounters mirror surface (globs normally avoid mirrors). The system predicted the glob might split or duplicate, creating linked mirror-globs through combining unusual properties. Most creative prediction, appropriately marked low confidence.

T5: Flutter seeds near micro-resonators (emit whisp-only frequencies). The system inferred frequency incompatibility precluding direct effects, but predicted indirect effects through whisp air currents, demonstrating multi-step causal reasoning. Confidence: Medium-High for no direct effect, Medium for indirect effects.

Results

The system successfully generated predictions for all five scenarios (5/5) with appropriate confidence calibration (5/5), proposing three novel mechanisms not present in original descriptions. Confidence levels aligned with prediction basis throughout.

Evidence for core mechanisms: Knowledge representation compressed from approximately 2,000-3,000 words to 1,075 words (65-70% reduction) while preserving causal structure. The system made predictions using only 4 episodes and 9 generalizations, referencing specific episodes rather than reconstructing from first principles. Nine generalizations enabled predictions across diverse scenarios through flexible combination. The system never claimed to "remember" scenarios – responses took the form "Based on generalizations G1 and G5, I predict..." Each prediction identified what to test next without external prompting.

Comparison with JEPA

Both approaches emphasize world models, prediction validation, sample efficiency, and hierarchical abstraction. Key differences:

Learning source: JEPA uses embodied interaction; episodic compression uses narrative-formatted experiences.

Latent variables: JEPA learns representations through gradient descent on sensory data; episodic compression creates episodes and generalizations through salience-based compression.

Training requirements: JEPA requires training from scratch over weeks/months; episodic compression leverages pre-trained models through architectural design, enabling immediate deployment.

These approaches may be complementary rather than contradictory. Embodied learning may prove optimal for physical manipulation tasks requiring sensorimotor grounding. Episodic compression may prove optimal for abstract reasoning and social intelligence where knowledge exists in narrative form.

Scaling Through Two-Layer Architecture

Part 1 validation occurred at "toy scale" – 9 generalizations and 4 episodes fitting comfortably within a single context window. Real-world deployment requires scaling to hundreds or thousands of generalizations.

Part 2 proposes functional separation into two components:

Intuition maintains consolidated knowledge and performs pattern-matching. It uses small efficient models (1-3B parameters) with large context windows (100K-200K tokens) for comprehensive pattern-matching, storing all generalizations and select high-salience episodes, running continuously in parallel with Executive processing.

Executive handles reasoning and consolidation. It uses large capable models with standard context windows for deliberate analysis, working with retrieved relevant subsets rather than full knowledge bases, conducting periodic consolidation during brief "micro-sleep" cycles, and engaging selectively when Intuition flags novelty or users initiate complex queries.

This separation exploits different computational properties: Intuition leverages transformer attention for parallel pattern-matching across extensive context; Executive performs serial depth-first reasoning within limited context.

Multi-Timescale Learning

Part 3 explores an emergent property: the architecture naturally operates across multiple timescales mirroring aspects of human learning. Real-time (seconds) involves Intuition pattern-matching against stored generalizations. Episodic formation (minutes-hours) involves salience detection identifying surprising events. Consolidation (hours-days) involves micro-sleep cycles extracting patterns. Long-term accumulation (weeks-months-years) involves generalization growth and refinement during deployment.

This temporal structure addresses a gap in current ML: continuous learning during deployment rather than distinct training/deployment phases. Traditional ML operates through training (weeks-months processing millions of examples), deployment (indefinite, model frozen), and occasional retraining (risks catastrophic forgetting). Episodic compression learns continuously during operation through frequent brief updates rather than infrequent massive retraining.

Practical Implications

The approach offers several advantages:

- Works with existing pre-trained models without training costs
- Enables immediate deployment and rapid prototyping
- Provides transparency through human-readable episodes/generalizations
- Allows cultural adaptability and cost-efficient operation
- Enables personalization without configuration – each deployment learns its specific context automatically

Systems improve over time in deployment: Week 1 performance < Month 1 < Month 6. Value increases with operational duration, creating switching costs from accumulated context-specific knowledge.

Limitations and Future Work

Current limitations require acknowledgment:

Unproven at scale: Validated at toy scale. Real-world deployment requires hundreds or thousands of generalizations.

No adaptation validation: Tested prediction but not learning from failures.

Uncertain thresholds: The proposed 3-5 episode threshold requires empirical validation.

Single-session only: Long-term stability over months or years remains unknown.

Artificial world limitations: Success in Shimmer Valleys doesn't guarantee real-world success.

Optimal parameters unknown: Salience thresholds, consolidation frequency, and other parameters require empirical tuning.

Future research must address real-world deployment across diverse domains for extended periods, threshold validation and parameter optimization, active learning and curiosity-driven exploration, hierarchical generalization and meta-patterns, multi-modal extension beyond language, and adversarial robustness testing.

Conclusion

This work suggests sample-efficient learning may not require embodied physical interaction. Humans learn extensively from narratives in practice. Large language models provide suitable cognitive substrates through pre-training and attention mechanisms. Episodic compression exploits these properties, achieving competence from minimal examples without neural network training.

Experimental validation demonstrates the mechanism works: 9 generalizations and 4 episodes enabled prediction in 5 novel scenarios with appropriate confidence calibration. The practical advantage is significant – this approach requires no new neural network training and enables immediate deployment with existing technology.

Whether episodic compression proves optimal or embodiment proves necessary for certain domains, recognizing that narrative learning provides a viable path forward has immediate practical value. The architecture enables systems that learn continuously during deployment, adapt to specific contexts automatically, and accumulate expertise without expensive retraining cycles – addressing critical challenges in deploying AI systems at scale.

References

Core Framework

1. LeCun, Y. (2022). A path towards autonomous machine intelligence. Open Review, 62.
2. Mossrake Group, LLC. (2025a). Learning Through Narratives: Episodic Compression to Latent Variables for Sample-Efficient Intelligence. Part 1 of a series.
3. Mossrake Group, LLC. (2025b). Scaling episodic compression: A two-layer architecture for continuous learning. Part 2 of a series.
4. Mossrake Group, LLC. (2025c). Multi-Timescale Continuous Learning: Emergent Properties of Episodic Compression. Part 3 of a series.

Episodic Memory and Learning

5. Tulving, E. (1972). Episodic and semantic memory. In E. Tulving & W. Donaldson (Eds.), Organization of memory (pp. 381-403). Academic Press.
6. Schacter, D. L., & Addis, D. R. (2007). The cognitive neuroscience of constructive memory: Remembering the past and imagining the future. Philosophical Transactions of the Royal Society B, 362(1481), 773-786.
7. Conway, M. A. (2009). Episodic memories. Neuropsychologia, 47(11), 2305-2313.
8. Ranganath, C., & Hsieh, L.-T. (2016). The hippocampus in time and space: Functional properties of episodic memory. Neuron, 90(2), 328–339.

Narrative Cognition

9. Bruner, J. (1986). Actual minds, possible worlds. Harvard University Press.
10. Bruner, J. (1990). Acts of meaning. Harvard University Press.
11. McAdams, D. P. (1993). The stories we live by: Personal myths and the making of the self. William Morrow.
12. Zacks, J. M., Speer, N. K., Swallow, K. M., Braver, T. S., & Reynolds, J. R. (2007). Event perception: A mind–brain perspective. Psychological Bulletin, 133(2), 273–293.
13. Kurby, C. A., & Zacks, J. M. (2008). Segmentation in narrative comprehension. Cognitive Psychology, 57(1), 25–61.

Curiosity and Intrinsic Motivation

14. Oudeyer, P. Y., & Kaplan, F. (2007). What is intrinsic motivation? A typology of computational approaches. Frontiers in Neurorobotics, 1, 6.
15. Schmidhuber, J. (2010). Formal theory of creativity, fun, and intrinsic motivation (1990-2010). IEEE Transactions on Autonomous Mental Development, 2(3), 230-247.

Sample Efficiency and Few-Shot Learning

16. Lake, B. M., Salakhutdinov, R., & Tenenbaum, J. B. (2015). Human-level concept learning through probabilistic program induction. Science, 350(6266), 1332-1338.
17. Lake, B. M., Ullman, T. D., Tenenbaum, J. B., & Gershman, S. J. (2017). Building machines that learn and think like people. Behavioral and Brain Sciences, 40, e253.

Causal Learning

18. Pearl, J. (2009). Causality: Models, reasoning, and inference (2nd ed.). Cambridge University Press.
19. Gopnik, A., & Schulz, L. (2007). Causal learning: Psychology, philosophy, and computation. Oxford University Press.
20. Tenenbaum, J. B., Lake, B. M., Kemp, C., & Gershman, S. J. (2023). Compositional causal learning and time. Cognitive Science, 47(2), e13209.

Value Learning

21. Russell, S. (2019). Human compatible: Artificial intelligence and the problem of control. Viking.
22. Christiano, P., Leike, J., Brown, T. B., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. Advances in Neural Information Processing Systems, 30.

World Models

23. Ha, D., & Schmidhuber, J. (2018). World models. arXiv preprint arXiv:1803.10122.
24. Hafner, D., Lillicrap, T., Ba, J., & Norouzi, M. (2019). Dream to control: Learning behaviors by latent imagination. arXiv preprint arXiv:1912.01603.

Memory-Augmented Networks

25. Graves, A., Wayne, G., & Danihelka, I. (2014). Neural Turing machines. arXiv preprint arXiv:1410.5401.
26. Graves, A., Wayne, G., Reynolds, M., et al. (2016). Hybrid computing using a neural network with dynamic external memory. Nature, 538(7626), 471-476.

Meta-Learning

27. Finn, C., Abbeel, P., & Levine, S. (2017). Model-agnostic meta-learning for fast adaptation of deep networks. International Conference on Machine Learning (ICML), 1126-1135.

Retrieval-Augmented Generation

28. Lewis, P., et al. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. Advances in Neural Information Processing Systems, 33, 9459-9474.

Cognitive Systems

29. Kahneman, D. (2011). Thinking, fast and slow. Farrar, Straus and Giroux.
30. Schneider, W., & Shiffrin, R. M. (1977). Controlled and automatic human information processing: I. Detection, search, and attention. Psychological Review, 84(1), 1-66.
31. Shiffrin, R. M., & Schneider, W. (1977). Controlled and automatic human information processing: II. Perceptual learning, automatic attending and a general theory. Psychological Review, 84(2), 127-190.

Catastrophic Forgetting

32. McCloskey, M., & Cohen, N. J. (1989). Catastrophic interference in connectionist networks: The sequential learning problem. Psychology of Learning and Motivation, 24, 109-165.
33. French, R. M. (1999). Catastrophic forgetting in connectionist networks. Trends in Cognitive Sciences, 3(4), 128-135.

Continual Learning

34. Kirkpatrick, J., et al. (2017). Overcoming catastrophic forgetting in neural networks. Proceedings of the National Academy of Sciences, 114(13), 3521-3526.
35. Rusu, A. A., et al. (2016). Progressive neural networks. arXiv preprint arXiv:1606.04671.

Cognitive Architecture

36. Schmidhuber, J. (2010). Formal theory of creativity, fun, and intrinsic motivation (1990–2010). IEEE Transactions on Autonomous Mental Development, 2(3), 230–247.
37. Ha, D., & Schmidhuber, J. (2018). World Models. arXiv:1803.10122.
38. Hafner, D., Lillicrap, T., Ba, J., Fischer, I., & van den Oord, A. (2019). Learning latent dynamics for planning from pixels. Proceedings of ICLR.
39. Graves, A., Wayne, G., & Danihelka, I. (2014). Neural Turing Machines. arXiv:1410.5401.
40. Graves, A., Wayne, G., Reynolds, M., Harley, T., et al. (2016). Hybrid computing using a neural network with dynamic external memory. Nature, 538, 471–476.

Memory and Learning

41. Tulving, E. (1972). Episodic and semantic memory. In E. Tulving & W. Donaldson (Eds.), Organization of Memory (pp. 381–403). Academic Press.
42. Conway, M. A. (2009). Episodic memories. Neuropsychologia, 47(11), 2305–2313.

43. Schacter, D. L., & Addis, D. R. (2007). Constructive memory: The ghosts of past and future. Nature, 445(7123), 27–30.
44. Ranganath, C., & Hsieh, L.-T. (2016). The hippocampus in time and space: Functional properties of episodic memory. Neuron, 90(2), 328–339.

Temporal Modeling and Multi-Timescale Learning

45. Dai, Z., Wang, Y., & Dong, L. (2024). Transformer architectures for event time modeling. Proceedings of ICLR 2024.
46. Ahuja, A., & Singh, A. (2023). Memory-augmented transformers for continual learning. Proceedings of ICML 2023.
47. Kazemi, S. M., Goel, R., Kipf, T., Kazemi, A., Brubaker, M., & Hamilton, W. L. (2023). Temporal knowledge graphs: A survey. Journal of Artificial Intelligence Research, 77, 1335–1385.
48. Botvinick, M., Wang, J. X., Dabney, W., Miller, K. J., & Kurth-Nelson, Z. (2019). Reinforcement learning, fast and slow. Trends in Cognitive Sciences, 23(5), 408–422.
49. Kiebel, S. J., Daunizeau, J., & Friston, K. J. (2008). A hierarchy of time-scales in the brain. PLoS Computational Biology, 4(11), e1000209.
50. Tallec, C., & Ollivier, Y. (2018). Can recurrent neural networks warp time? Proceedings of ICLR 2018.

Version 1.0